

Persian Gulf University
Deep Learning
Spring Semester 2022-2023

Teacher: Dr. Rostami

Ali Behrouzi

Stdn: 4010724002

Email: alibehrouzi97@gmail.com

Assignment 6

- CNN , cat and dog dataset

I used a 200 mb database , include training and test folders, each include dog and cat folders. It has been drive from kaggle dataset.

At the end I ended up with 2 piece of code , first doesn't include dropout and just have a slight augmentation , but the second includes dropout layer and more complex augmentation.

I uploaded my dataset to google drive so I needed to mount my drive to access my dataset :

Convolutional Neural Network

```
[ ] from google.colab import drive  
    drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

Then import libraries :

Importing the libraries

```
[ ] import tensorflow as tf  
    from keras.preprocessing.image import ImageDataGenerator
```

Then we should preprocess the data :

In first code A : I just used normalization and 3 option for data augmentation :

Part 1 - Data Preprocessing

Preprocessing the Training set

```
▶ train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True  
)
```

But in code B : I added 4 other augmentation option to the code and the rest is the same :

▼ Part 1 - Data Preprocessing

▼ Preprocessing the Training set

```
▶ train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest'
```

```
)
```

The we have to tell the program from which source it should get the training images, and also the batch size and if the class mode is binary or categorical:

```
training_set= train_datagen.flow_from_directory(  
    '/content/gdrive/My Drive/dataset/training_set',  
    target_size=(150,150),  
    batch_size=32,  
    class_mode='binary'  
)
```

Found 8000 images belonging to 2 classes.

Do the same preprocessing for test data but it just need to be normalized :

▼ Preprocessing the Test set

```
test_datagen = ImageDataGenerator(  
    rescale=1./255)  
test_set=test_datagen.flow_from_directory(  
    '/content/gdrive/My Drive/dataset/test_set',  
    target_size=(150,150),  
    batch_size=32,  
    class_mode='binary'  
)
```

Found 2000 images belonging to 2 classes.

Start building the cnn : create a sequential model and add a convolutional layer to it the apply maxpooling . you should tell it how many filters you want, and the size of filters , activation function and input size of this layer .stride is (1,1) by default. the size of window for maxpooling should be given to it :

▼ Part 2 - Building the CNN

▼ Initialising the CNN

```
[ ] cnn =tf.keras.models.Sequential()
```

▼ Step 1 - Convolution

```
[ ] cnn.add(tf.keras.layers.Conv2D(filters=32,kernel_size=3,activation='relu',input_shape=[150,150,3]))
```

▼ Step 2 - Pooling

```
[ ] cnn.add(tf.keras.layers.MaxPooling2D((2,2)))
```

We add 3 more conv layers :

▼ Adding three more convolutional layers

```
[ ] cnn.add(tf.keras.layers.Conv2D(filters=64,kernel_size=3,activation='relu'))  
cnn.add(tf.keras.layers.MaxPooling2D((2,2)))
```

```
[ ] cnn.add(tf.keras.layers.Conv2D(filters=128,kernel_size=3,activation='relu'))  
cnn.add(tf.keras.layers.MaxPooling2D((2,2)))
```

```
[ ] cnn.add(tf.keras.layers.Conv2D(filters=128,kernel_size=3,activation='relu'))  
cnn.add(tf.keras.layers.MaxPooling2D((2,2)))
```

Then before giving the output of conv layers to dense layer, you should flatten the result :

▼ Step 3 - Flattening

```
[ ] cnn.add(tf.keras.layers.Flatten())
```

Then we should add dense layer :

In code A we didn't use dropout layer :

▼ Step 4 - Full Connection

```
[ ] cnn.add(tf.keras.layers.Dense(units=512,activation='relu'))
```

▼ Step 5 - Output Layer

```
[ ] cnn.add(tf.keras.layers.Dense(units=1,activation='sigmoid'))
```

But in code B , we added dropout layer for reducing overfitting . we used sigmoid for last layer because it's just binary classification:

▼ Step 4 - Full Connection

```
[ ] cnn.add(tf.keras.layers.Dropout(0.5))
```

```
[ ] cnn.add(tf.keras.layers.Dense(units=512,activation='relu'))
```

▼ Step 5 - Output Layer

```
[ ] cnn.add(tf.keras.layers.Dense(units=1,activation='sigmoid'))
```

Then in training stage we first compile the model with adam optimizer and binary_crossentropy loss function and then call fit function with 30 epochs :

▼ Part 3 - Training the CNN

▼ Compiling the CNN

```
[ ] cnn.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
```

▼ Training the CNN on the Training set and evaluating it on the Test set

```
▶ history=cnn.fit(x=training_set,validation_data=test_set,epochs=30)
```

This is the result for code A from epoch 17, it seems from epoch 16, the val_acc doesn't change significantly but train_acc increases (overfitting) :

```
Epoch 16/30
250/250 [=====] - 88s 351ms/step - loss: 0.2179 - accuracy: 0.9051 - val_loss: 0.2857 - val_accuracy: 0.8890
Epoch 17/30
250/250 [=====] - 90s 359ms/step - loss: 0.2030 - accuracy: 0.9145 - val_loss: 0.3195 - val_accuracy: 0.8760
Epoch 18/30
250/250 [=====] - 88s 351ms/step - loss: 0.1919 - accuracy: 0.9178 - val_loss: 0.3184 - val_accuracy: 0.8780
Epoch 19/30
250/250 [=====] - 87s 347ms/step - loss: 0.1842 - accuracy: 0.9246 - val_loss: 0.2958 - val_accuracy: 0.8880
Epoch 20/30
250/250 [=====] - 90s 358ms/step - loss: 0.1667 - accuracy: 0.9340 - val_loss: 0.3244 - val_accuracy: 0.8865
Epoch 21/30
250/250 [=====] - 87s 349ms/step - loss: 0.1618 - accuracy: 0.9335 - val_loss: 0.3619 - val_accuracy: 0.8705
Epoch 22/30
250/250 [=====] - 89s 355ms/step - loss: 0.1382 - accuracy: 0.9449 - val_loss: 0.3553 - val_accuracy: 0.8840
Epoch 23/30
250/250 [=====] - 90s 362ms/step - loss: 0.1302 - accuracy: 0.9488 - val_loss: 0.3663 - val_accuracy: 0.8915
Epoch 24/30
250/250 [=====] - 86s 346ms/step - loss: 0.1191 - accuracy: 0.9534 - val_loss: 0.3383 - val_accuracy: 0.8850
Epoch 25/30
250/250 [=====] - 87s 350ms/step - loss: 0.1162 - accuracy: 0.9559 - val_loss: 0.3703 - val_accuracy: 0.8820
Epoch 26/30
250/250 [=====] - 89s 355ms/step - loss: 0.1049 - accuracy: 0.9607 - val_loss: 0.3567 - val_accuracy: 0.8985
Epoch 27/30
250/250 [=====] - 92s 368ms/step - loss: 0.0985 - accuracy: 0.9605 - val_loss: 0.4422 - val_accuracy: 0.8810
Epoch 28/30
250/250 [=====] - 89s 358ms/step - loss: 0.0987 - accuracy: 0.9619 - val_loss: 0.3861 - val_accuracy: 0.8950
Epoch 29/30
250/250 [=====] - 90s 360ms/step - loss: 0.0911 - accuracy: 0.9672 - val_loss: 0.4002 - val_accuracy: 0.8880
Epoch 30/30
250/250 [=====] - 87s 349ms/step - loss: 0.0862 - accuracy: 0.9686 - val_loss: 0.4878 - val_accuracy: 0.8750
```

And this is the result of code B from epoch 16 :

```
Epoch 16/30
250/250 [=====] - 672s 3s/step - loss: 0.4952 - accuracy: 0.7586 - val_loss: 0.4114 - val_accuracy: 0.8145
Epoch 17/30
250/250 [=====] - 706s 3s/step - loss: 0.4806 - accuracy: 0.7755 - val_loss: 0.4154 - val_accuracy: 0.8120
Epoch 18/30
250/250 [=====] - 695s 3s/step - loss: 0.4782 - accuracy: 0.7707 - val_loss: 0.3855 - val_accuracy: 0.8260
Epoch 19/30
250/250 [=====] - 699s 3s/step - loss: 0.4658 - accuracy: 0.7824 - val_loss: 0.3559 - val_accuracy: 0.8520
Epoch 20/30
250/250 [=====] - 660s 3s/step - loss: 0.4446 - accuracy: 0.7904 - val_loss: 0.3768 - val_accuracy: 0.8340
Epoch 21/30
250/250 [=====] - 672s 3s/step - loss: 0.4349 - accuracy: 0.7991 - val_loss: 0.3481 - val_accuracy: 0.8540
Epoch 22/30
250/250 [=====] - 709s 3s/step - loss: 0.4352 - accuracy: 0.7991 - val_loss: 0.3348 - val_accuracy: 0.8575
Epoch 23/30
250/250 [=====] - 658s 3s/step - loss: 0.4220 - accuracy: 0.8062 - val_loss: 0.3797 - val_accuracy: 0.8275
Epoch 24/30
250/250 [=====] - 694s 3s/step - loss: 0.4228 - accuracy: 0.8050 - val_loss: 0.3277 - val_accuracy: 0.8570
Epoch 25/30
250/250 [=====] - 677s 3s/step - loss: 0.4059 - accuracy: 0.8141 - val_loss: 0.3135 - val_accuracy: 0.8680
Epoch 26/30
250/250 [=====] - 624s 2s/step - loss: 0.4070 - accuracy: 0.8175 - val_loss: 0.3105 - val_accuracy: 0.8620
Epoch 27/30
250/250 [=====] - 670s 3s/step - loss: 0.3967 - accuracy: 0.8164 - val_loss: 0.3377 - val_accuracy: 0.8490
Epoch 28/30
250/250 [=====] - 680s 3s/step - loss: 0.3925 - accuracy: 0.8220 - val_loss: 0.3170 - val_accuracy: 0.8745
Epoch 29/30
250/250 [=====] - 699s 3s/step - loss: 0.3813 - accuracy: 0.8290 - val_loss: 0.3301 - val_accuracy: 0.8510
Epoch 30/30
250/250 [=====] - 662s 3s/step - loss: 0.3776 - accuracy: 0.8296 - val_loss: 0.2951 - val_accuracy: 0.8720
Epoch 30/30
250/250 [=====] - 635s 3s/step - loss: 0.3696 - accuracy: 0.8321 - val_loss: 0.3148 - val_accuracy: 0.8675
```


Then I tried to plot the result in all epochs :

```
import matplotlib.pyplot as plt

history_dict = history.history


epochs = range(1, 30 + 1)

acc_values = history_dict['accuracy']
val_acc_values = history_dict['val_accuracy']


plt.plot(epochs, acc_values, 'bo', label='Training acc')
plt.plot(epochs, val_acc_values, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

This is the result for code A , it seems from epoch 16 the model start overfitting :



This is the result for code B :



Then I tried to find the best epoch(based on val_acc) and val_acc and train_acc of that epoch :

This is the result for code A :

```
[ ] # Find the epoch with the best validation accuracy
    best_epoch = np.argmax(val_acc_values)

    # Print the training and validation accuracy for the best epoch
    print(f"Best Epoch: {best_epoch + 1}")
    print(f"Training Accuracy: {acc_values[best_epoch]}")
    print(f"Validation Accuracy: {val_acc_values[best_epoch]}")
```

```
Best Epoch: 26
Training Accuracy: 0.9607499837875366
Validation Accuracy: 0.8985000252723694
```

This is the result for code B :

```
0s # Find the epoch with the best validation accuracy
best_epoch = np.argmax(val_acc_values)

# Print the training and validation accuracy for the best epoch
print(f"Best Epoch: {best_epoch + 1}")
print(f"Training Accuracy: {acc_values[best_epoch]}")
print(f"Validation Accuracy: {val_acc_values[best_epoch]}")

Best Epoch: 27
Training Accuracy: 0.8220000267028809
Validation Accuracy: 0.8744999766349792
```

Then I just try to predict one single image, and the model predicted correctly :

```
Part 4 - Making a single prediction

from tensorflow.keras.preprocessing import image
test_image = image.load_img('/content/gdrive/My Drive/dataset/single_prediction/cat_or_dog_1.jpg', target_size = (150, 150)) # this is a dog picture
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis = 0)
result = cnn.predict(test_image)
training_set.class_indices
if result[0][0] == 1:
    prediction = 'dog'
else:
    prediction = 'cat'

print(prediction)

1/1 [=====] - 0s 279ms/step
dog
```